

## Lecture 13

*Course Coordinator: Prof. Fei-Fei Li**Scribes: Akash Gupta*

**Unsupervised learning** - Unlabeled training data and the goal is to learn some underlying structure of data. E.g. - Clustering, Dimensionality Reduction, density estimation, etc.

**Generative Models** - Given training data, generate new samples from same distribution.

Why care? -

- Generate realistic samples, super-resolution, etc.
- Generative modeling of time-series data can be used for simulation and planning.
- Training generative models also enable inference of latent representations that can be useful as general features.

– > Taxonomy of Generative Models - Ian Goodfellow. Tutorial on GANs. 2017

**Pixel RNN & Pixel CNN** -

- Fully Visible Belief Networks.
- Explicit density model.
- Use chain rule to decompose likelihood of an image  $x$  into product of 1-d distributions and then maximize likelihood of the training data.
- Chain rule allows to define a tractable density function.

$$p(x) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

Since it's a complex distribution, use neural networks to model this.

**Pixel RNN** -

- Generate image pixels starting from corner.
- Dependency on previous pixels modeled using an RNN(LSTM).
- sequential generation is slow.

**PixelCNN** -

- Still generate pixels starting from the corner.
- Dependency of previous pixels now modeled using pixels in the context region.
- Training is faster than PixelRNN. Generation time still slow since sequential generation.

## Variational Autoencoders -

Background: Autoencoders - Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data. Useful for dimensionality reduction.

$$x \rightarrow z$$

$z$  latent variable,  $x$  feature variable

- > Train such that latent variables can be used to reconstruct original data - "Autoencoding".
- > Use L2 loss for training.
- > Throw away the decoder and fine tune the encoder as a classifier for supervised task.

Q) Why  $z$  should be smaller in dimension than  $x$ ?

Ans)  $z$  should represent the most important features in  $x$  OR  $z$  should capture meaningful factors of variation.

VAEs - Probabilistic spin on traditional autoencoders.

- Assume training data is generated from some underlying latent representation  $z$ .
- Representing the model -
  - Choose prior  $p(z)$  to be simple. e.g. - Gaussian
  - Conditional  $p(x|z)$  is complex (generates image) - > represent with NN - Decoder network.
- Training the model -
  - Maximize the data likelihood - (In this case prior and conditional are continuous functions)

$$p_{\theta}(x) = \int p_{\theta}(z)p_{\theta}(x|z)dz$$

- Problem - Intractable function. Posterior also tractable.

$$p_{\theta}(z|x) = p_{\theta}(x|z)p_{\theta}(z)/p_{\theta}(x)$$

- Solution - Estimate posterior using an encoder network that defines  $q_{\phi}(z|x)$  and approximates  $p_{\theta}(z|x)$  - Allows to derive a lower bound on data likelihood that is tractable and we can optimize.
- Encoder Network (also called recognition/inference network) - Gives mean and covariance of  $z|x$  and sample from this.
- Decoder Network (also called generation network) - Gives mean and covariance of  $x|z$  and sample from this.
- Maximizing the likelihood lower bound (also called ELBO - Evidence lower bound, since it's a value  $\geq \log$  of data likelihood. This is done in order to make the data likelihood a tractable function) -

$$E_z[\log p_{\theta}(x^{(i)}|z)] - D_{KL}(q_{\phi}(z|x^{(i)})||p_{\theta}(z))$$

- first term gives the reconstruction.
- second term is minimizing the distance approximate dist. and prior (Gaussian in this case)
- For generation, just use the decoder on random samples drawn from prior to output new data. (Cons - Produces blurry images)

**Generative Adversarial Networks (GANs)** - Implicit density estimation. A game-theoretic approach.

- Problem - To generate data, sample from high dimensional training distribution. No direct way to do this.
- Sample from a simple distribution, e.g. random noise. Learn the transformation using a NN to a training distribution.
- Represent model - 2-player game -
  - Generator network - try to fool the discriminator by generating real looking images.
  - Discriminator network - try to distinguish between real and fake images.
- Training -
  - Train jointly as minimax game.

$$\min_{\theta_g} \max_{\theta_d} [E_{x \sim p_{data}} \log D_{\theta_d}(x) + E_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))]$$

- Discriminator ( $\theta_d$ ) wants to maximize objective such that  $D(x)$  is close to 1(real) and  $D(G(z))$  is close to 0(fake).
- Generator ( $\theta_g$ ) wants to minimize objective such that  $D(G(z))$  is close to 1(trying to fool the discriminator)
- Gradient Ascent to maximize the objective for discriminator.
- Gradient Ascent to maximize the flipped objective for generator since real objective is steep when samples are good but flat when samples are bad (small gradients). Flipping reverses this trend.
- Algorithm - For  $k$  steps, train only the discriminator. After that update generator.
- $k = 1$  or  $k > 1$ , no best rule. E.g. - Wasserstein GAN alleviates the problem, better stability.
- After training, use generator to generate new images.
- E.g. Unsupervised Representation learning with DCGAN. Radford et. al. ICLR. 2016, The GAN Zoo.
- Tips and tricks for training GANs - [Link](#)
- Beautiful, state-of-the-art samples generated but trickier to train